

DA5030.A3.Parpattedar

Shruti Parpattedar

February 3, 2019

Question 1

Downloading and loading the dataset into R

```
#setwd("D:/NEU/DA5030/Assignment3")
prc <- read.csv("prostate_cancer.csv", stringsAsFactors = FALSE)
```

Question 2

Preparing and exploring the data

```
str(prc)

## 'data.frame':    100 obs. of  10 variables:
## $ id          : int  1 2 3 4 5 6 7 8 9 10 ...
## $ diagnosis_result : chr  "M" "B" "M" "M" ...
## $ radius       : int  23 9 21 14 9 25 16 15 19 25 ...
## $ texture      : int  12 13 27 16 19 25 26 18 24 11 ...
## $ perimeter    : int  151 133 130 78 135 83 120 90 88 84 ...
## $ area         : int  954 1326 1203 386 1297 477 1040 578 520 476 ...
## $ smoothness   : num  0.143 0.143 0.125 0.07 0.141 0.128 0.095 0.119 0.127 0.119 ...
## $ compactness  : num  0.278 0.079 0.16 0.284 0.133 0.17 0.109 0.165 0.193 0.24 ...
## $ symmetry     : num  0.242 0.181 0.207 0.26 0.181 0.209 0.179 0.22 0.235 0.203 ...
## $ fractal_dimension: num  0.079 0.057 0.06 0.097 0.059 0.076 0.057 0.075 0.074 0.082 ...

prc <- prc[-1]
table(prc$diagnosis_result)

##
##  B  M
## 38 62

prc$diagnosis <- factor(prc$diagnosis_result, levels = c("B", "M"),
                        labels = c("Benign", "Malignant"))
round(prop.table(table(prc$diagnosis)) * 100, digits = 1)

##
##      Benign Malignant
##          38         62

Normalizing numeric data

normalize <- function(x) {
  return ((x - min(x)) / (max(x) - min(x))) }
prc_n <- as.data.frame(lapply(prc[2:9], normalize))
summary(prc_n$radius)

##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## 0.0000  0.1875  0.5000  0.4906  0.7500  1.0000
```

Creating training and test data set

```
prc_train <- prc_n[1:65,]  
prc_test <- prc_n[66:100,]  
prc_train_labels <- prc[1:65, 1]  
prc_test_labels <- prc[66:100, 1]
```

Training a model on data

```
#install.packages("class")  
library(class)  
prc_test_pred <- knn(train = prc_train, test = prc_test, cl = prc_train_labels, k=10)
```

Evaluate the model performance

Accuracy - ((TN+TP)/35) = 0.63%

```
#install.packages("gmodels")  
library(gmodels)  
CrossTable(prc_test_labels, prc_test_pred, prop.chisq = FALSE)
```

```
##  
##  
##      Cell Contents  
## |-----|  
## |                      N |  
## |      N / Row Total |  
## |      N / Col Total |  
## |      N / Table Total |  
## |-----|  
##  
##  
## Total Observations in Table:  35  
##  
##  
##      | prc_test_pred  
## prc_test_labels |      B |      M | Row Total |  
## -----|-----|-----|-----|  
##      B |      6 |     13 |      19 |  
##      |    0.316 |    0.684 |    0.543 |  
##      |    1.000 |    0.448 |          |  
##      |    0.171 |    0.371 |          |  
## -----|-----|-----|-----|  
##      M |      0 |     16 |      16 |  
##      |    0.000 |    1.000 |    0.457 |  
##      |    0.000 |    0.552 |          |  
##      |    0.000 |    0.457 |          |  
## -----|-----|-----|-----|  
##      Column Total |      6 |     29 |      35 |  
##      |    0.171 |    0.829 |          |  
## -----|-----|-----|-----|  
##  
##
```

Improve the performance of the model

Using k=9, I am getting 0 false negatives which is an improvement over 1 false negative which was being observed k=10.

Accuracy - $((TN+TP)/35) = 0.69\%$

Using $k=11$, I am getting 0 false negatives but 14 false positives.

Accuracy - $((TN+TP)/35) = 0.6\%$

```
prc_test_pred2 <- knn(train = prc_train, test = prc_test, cl = prc_train_labels, k=9)
CrossTable(prc_test_labels, prc_test_pred2, prop.chisq = FALSE)
```

```
##
##
##      Cell Contents
## |-----|
## |                N |
## |      N / Row Total |
## |      N / Col Total |
## |      N / Table Total |
## |-----|
##
##
## Total Observations in Table:  35
##
##
##      | prc_test_pred2
## prc_test_labels |          B |          M | Row Total |
## -----|-----|-----|-----|
##          B |          8 |          11 |          19 |
##          |      0.421 |      0.579 |      0.543 |
##          |      1.000 |      0.407 |          |
##          |      0.229 |      0.314 |          |
## -----|-----|-----|-----|
##          M |          0 |          16 |          16 |
##          |      0.000 |      1.000 |      0.457 |
##          |      0.000 |      0.593 |          |
##          |      0.000 |      0.457 |          |
## -----|-----|-----|-----|
##      Column Total |          8 |          27 |          35 |
##          |      0.229 |      0.771 |          |
## -----|-----|-----|-----|
##
##
```

```
prc_test_pred3 <- knn(train = prc_train, test = prc_test, cl = prc_train_labels, k=11)
CrossTable(prc_test_labels, prc_test_pred3, prop.chisq = FALSE)
```

```
##
##
##      Cell Contents
## |-----|
## |                N |
## |      N / Row Total |
## |      N / Col Total |
## |      N / Table Total |
## |-----|
##
##
```

```
## Total Observations in Table: 35
##
##
##      | prc_test_pred3
## prc_test_labels |      B |      M | Row Total |
## -----|-----|-----|-----|
##      B |      5 |     14 |      19 |
##      |    0.263 |    0.737 |    0.543 |
##      |    1.000 |    0.467 |      |
##      |    0.143 |    0.400 |      |
## -----|-----|-----|-----|
##      M |      0 |     16 |      16 |
##      |    0.000 |    1.000 |    0.457 |
##      |    0.000 |    0.533 |      |
##      |    0.000 |    0.457 |      |
## -----|-----|-----|-----|
## Column Total |      5 |     30 |      35 |
##      |    0.143 |    0.857 |      |
## -----|-----|-----|-----|
##
##
```

Question 3

Using the kNN algorithm from the caret package.

```
library(caret)
```

```
## Loading required package: lattice
```

```
## Loading required package: ggplot2
```

```
library(doSNOW)
```

```
## Loading required package: foreach
```

```
## Loading required package: iterators
```

```
## Loading required package: snow
```

```
library(xgboost)
```

```
# Loading the data onto a different variable in R and removing the id column from it.
```

```
data <- read.csv("prostate_cancer.csv", stringsAsFactors = FALSE)
```

```
#str(data)
```

```
data <- data[,-1]
```

```
# Using the diagnosis_result column as a factor instead of a plain character.
```

```
data$diagnosis_result <- as.factor(data$diagnosis_result)
```

```
# Partitioning the data into a 65-35 training and testing dataset.
```

```
set.seed(300)
```

```
indexes <- createDataPartition(data$diagnosis_result, p = 0.64, list = FALSE)
```

```
data.train <- data[indexes,]
```

```
data.test <- data[-indexes,]
```

```
#prop.table(table(data$diagnosis_result))
```

```
#prop.table(table(data.train$diagnosis_result))
```

```

#prop.table(table(data.test$diagnosis_result))

trainX <- data.train[,names(data.train) != "diagnosis_result"]
preProcValues <- preProcess(x = trainX, method = c("center","scale"))

# Training and training control from the dataset.
set.seed(400)
train.control <- trainControl(method = "repeatedcv", repeats = 3)

# Finding the knn fit for the training set and then plotting it.
knnFit <- train(diagnosis_result ~ .,
               data = data.train,
               method = "knn",
               trControl = train.control,
               preProcess = c("center","scale"),
               tuneLength = 20)
knnFit


```

```

## k-Nearest Neighbors
##
## 65 samples
## 8 predictor
## 2 classes: 'B', 'M'
##
## Pre-processing: centered (8), scaled (8)
## Resampling: Cross-Validated (10 fold, repeated 3 times)
## Summary of sample sizes: 59, 58, 58, 59, 59, 58, ...
## Resampling results across tuning parameters:
##
##  k  Accuracy  Kappa
##  5  0.7896825  0.53348190
##  7  0.8103175  0.56985871
##  9  0.8214286  0.59207886
## 11  0.8317460  0.60709267
## 13  0.8492063  0.64939839
## 15  0.8373016  0.62536458
## 17  0.8222222  0.59572432
## 19  0.8158730  0.57391273
## 21  0.8412698  0.62402754
## 23  0.8515873  0.65318627
## 25  0.8634921  0.67637468
## 27  0.8261905  0.57681912
## 29  0.8007937  0.51057155
## 31  0.8111111  0.53963235
## 33  0.8134921  0.54240542
## 35  0.7936508  0.48910220
## 37  0.7333333  0.33073593
## 39  0.6849206  0.18744589
## 41  0.6412698  0.07142857
## 43  0.6190476  0.00000000
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was k = 25.

```

```
plot(knnFit)
```



```
# Using the model to predict data on the testing set.
```

```
knnPredict <- predict(knnFit, newdata = data.test)
```

```
knnPredict
```

```
## [1] M M M M M M M M B M B M M M B M M M B B M M B M M M M M M M B M B
## Levels: B M
```

Question 4

Generating confusion matrices for the kNN predictions made using the two algorithms above.

```
prc_test_labels <- as.factor(prc_test_labels)
confusionMatrix(prc_test_pred, prc_test_labels)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  B  M
##           B   6  0
##           M  13 16
##
##           Accuracy : 0.6286
##           95% CI : (0.4492, 0.7853)
##           No Information Rate : 0.5429
##           P-Value [Acc > NIR] : 0.1987130
```

```
##
##           Kappa : 0.2968
## Mcnemar's Test P-Value : 0.0008741
##
##           Sensitivity : 0.3158
##           Specificity : 1.0000
##           Pos Pred Value : 1.0000
##           Neg Pred Value : 0.5517
##           Prevalence : 0.5429
##           Detection Rate : 0.1714
##           Detection Prevalence : 0.1714
##           Balanced Accuracy : 0.6579
##
##           'Positive' Class : B
##
confusionMatrix(knnPredict, data.test$diagnosis_result)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  B  M
##           B  8  0
##           M  5 22
##
##           Accuracy : 0.8571
##           95% CI : (0.6974, 0.9519)
##           No Information Rate : 0.6286
##           P-Value [Acc > NIR] : 0.002746
##
##           Kappa : 0.6679
## Mcnemar's Test P-Value : 0.073638
##
##           Sensitivity : 0.6154
##           Specificity : 1.0000
##           Pos Pred Value : 1.0000
##           Neg Pred Value : 0.8148
##           Prevalence : 0.3714
##           Detection Rate : 0.2286
##           Detection Prevalence : 0.2286
##           Balanced Accuracy : 0.8077
##
##           'Positive' Class : B
##
```